

計算機視覺作業

干皓丞，2101212850, 信息工程學院

2021 年 11 月 1 日

1 作業目標與章節摘要

根據 W6_MNIST_FC.ipynb 改良程式碼好獲得更好的效果。在此可以使用增加卷積層結構，或者嘗試增加 dropout 或者 BN 技術等，訓練出盡可能高的 MNIST 分類效果，原始程式碼在名為 kancheng/kancs-report-in-2021 的 Github 專案下，可以在 CV/pytorch-tensorflow-mnist/code 下找到。該報告第一節為 Pytorch MNIST Training 課堂程式碼改良，第二節為 TensorFlow MNIST Training，第三節為 CNN 方法與重要文獻，而附件則為 Pytorch MNIST Training 中的測試結果。

本次作业



- 在 W6_MNIST_FC.ipynb 基础上，增加卷积层结构/增加 dropout 或者 BN 技术等，训练出尽可能高的 MNIST 分类效果。



Fig. 1. 作業目標

2 Pytorch MNIST Training

Pytorch MNIST Training 原始的範例程式碼中 Test Accuracy 結果為 0.914，而自己 MacBook Pro (Retina, 15-inch, Mid 2014) 實際測試後，在預設值 EPOCH = 2 與 LR = 0.001 所跑出的 Test Accuracy 結果為 0.894 and 0.901。原始程式碼範例如下所示：

```
1 import torch
2 import torch.nn as nn
3 import torch.utils.data as Data
4 import torchvision
5 import torch.nn.functional as F
6 import numpy as np
7 # torch.manual_seed(1)
8 EPOCH = 2
9 LR = 0.001
10 DOWNLOAD_MNIST = True
11
12 train_data = torchvision.datasets.MNIST(root='./mnist/', train=True,
13                                         transform=torchvision.transforms.ToTensor(),
14                                         download=DOWNLOAD_MNIST, )
15 test_data = torchvision.datasets.MNIST(root='./mnist/', train=False)
16 train_x = torch.unsqueeze(train_data.data, dim=1).type(torch.FloatTensor)
17         / 255.
18 train_y = train_data.targets
19 print(train_x.shape)
20 test_x = torch.unsqueeze(test_data.data, dim=1).type(torch.FloatTensor)
21         [:2000] / 255. # Tensor on GPU
22 test_y = test_data.targets[:2000]
23 test_x.shape
24 import matplotlib.pyplot as plt
25 plt.imshow(test_x[1,0,:,:].numpy(), 'gray')
26 test_y[:10]
27 class FC(nn.Module):
28     def __init__(self):
29         super(FC, self).__init__()
30         self.fc1 = nn.Linear(784, 256)
31         self.fc2 = nn.Linear(256, 10)
32         self.fc3 = nn.Linear(10, 10)
33
34     def forward(self, x):
35         x = x.view(x.size(0), -1)
36         x = self.fc1(x)
```

```

34         x = F.relu(x)
35         x = self.fc2(x)
36         x = F.relu(x)
37         x = self.fc3(x)
38
39         output = x
40         return output
41
42
43 fc = FC()
44
45 optimizer = torch.optim.Adam(fc.parameters(), lr=LR)
46 # loss_func = nn.MSELoss()
47 loss_func = nn.CrossEntropyLoss()
48
49 data_size = 20000
50 batch_size = 50
51
52 for epoch in range(EPOCH):
53     random_indx = np.random.permutation(data_size)
54     for batch_i in range(data_size // batch_size):
55         indx = random_indx[batch_i * batch_size:(batch_i + 1) *
56             batch_size]
57
58         b_x = train_x[indx, :]
59         b_y = train_y[indx]
60         # print(b_x.shape)
61         # print(b_y.shape)
62
63         output = fc(b_x)
64         loss = loss_func(output, b_y)
65
66         loss.backward()
67         optimizer.step()
68         optimizer.zero_grad()
69
70     if batch_i % 50 == 0:
71         test_output = fc(test_x)
72         pred_y = torch.max(test_output, 1)[1].data.squeeze()
73         # pred_y = torch.max(test_output, 1)[1].data.squeeze()
74         accuracy = torch.sum(pred_y == test_y).type(torch.FloatTensor)

```

```

        ) / test_y.size(0)
74     print('Epoch: ', epoch, '| train loss: %.4f' % loss.data.cpu
           ().numpy(), '| test accuracy: %.3f' % accuracy)
75
76     test_output = fc(test_x[:10])
77     pred_y = torch.max(test_output, 1)[1].data.squeeze() # move the
           computation in GPU
78
79     print(pred_y, 'prediction number')
80     print(test_y[:10], 'real number')
81
82     test_output = fc(test_x[:1])
83     pred_y = torch.max(test_output, 1)[1].data.squeeze() # move the
           computation in GPU
84
85     print(pred_y, 'prediction number')
86     print(test_y[:1], 'real number')
87     test_output
88     test_x[:1].shape
89     plt.imshow(test_x[:1].numpy().squeeze(), 'gray')
90
91     import torch
92     torch.eye(10)
93     0, 1, 2, ..., 9

```

在此針對 FC class 和 batch_size 進行測試，一開始的 FC 則稱之為 FC3。FC 為控制訓練的層，而所謂的 Batch Size 為一個重要參數，該參數的意義在於做出在記憶體效率跟記憶體容量之間尋找最佳平衡，若編號為 FC3 且 Batch Size 為 50 則編號表述為 FC3-batch-50。此次實驗共有 FC3、FC4-1、FC4-2、FC7，Batch Size 為 10、30、50。

```

1  # FC3
2  class FC(nn.Module):
3      def __init__(self):
4          super(FC, self).__init__()
5          self.fc1 = nn.Linear(784, 256)
6          self.fc2 = nn.Linear(256, 10)
7          self.fc3 = nn.Linear(10, 10)
8
9      def forward(self, x):
10         x = x.view(x.size(0), -1)
11         x = self.fc1(x)
12         x = F.relu(x)

```

```

13         x = self.fc2(x)
14         x = F.relu(x)
15         x = self.fc3(x)
16         output = x
17         return output
18 fc = FC()

```

此 FC 則稱之為 FC4-1。

```

1 # FC4-1
2 class FC(nn.Module):
3     def __init__(self):
4         super(FC, self).__init__()
5         self.fc1 = nn.Linear(784, 256)
6         self.fc2 = nn.Linear(256, 128)
7         self.fc3 = nn.Linear(128, 10)
8         self.fc4 = nn.Linear(10, 10)
9
10    def forward(self, x):
11        x = x.view(x.size(0), -1)
12        x = self.fc1(x)
13        x = F.relu(x)
14        x = self.fc2(x)
15        x = F.relu(x)
16        x = self.fc3(x)
17        x = F.relu(x)
18        x = self.fc4(x)
19        x = F.relu(x)
20        output = x
21        return output
22 fc = FC()

```

此 FC 則稱之為 FC4-2。

```

1 # FC4-2
2 class FC(nn.Module):
3     def __init__(self):
4         super(FC, self).__init__()
5         self.fc1 = nn.Linear(784, 1256)
6         self.fc2 = nn.Linear(1256, 128)
7         self.fc3 = nn.Linear(128, 10)
8         self.fc4 = nn.Linear(10, 10)
9
10    def forward(self, x):

```

```

11         x = x.view(x.size(0), -1)
12         x = self.fc1(x)
13         x = F.relu(x)
14         x = self.fc2(x)
15         x = F.relu(x)
16         x = self.fc3(x)
17         x = F.relu(x)
18         x = self.fc4(x)
19         x = F.relu(x)
20         output = x
21         return output
22 fc = FC()

```

此 FC 則稱之為 FC7。

```

1 # FC7
2 class FC(nn.Module):
3     def __init__(self):
4         super(FC, self).__init__()
5         self.fc1 = nn.Linear(784, 256)
6         self.fc2 = nn.Linear(256, 128)
7         self.fc3 = nn.Linear(128, 64)
8         self.fc4 = nn.Linear(64, 32)
9         self.fc5 = nn.Linear(32, 16)
10        self.fc6 = nn.Linear(16, 10)
11        self.fc7 = nn.Linear(10, 10)
12
13    def forward(self, x):
14        x = x.view(x.size(0), -1)
15        x = self.fc1(x)
16        x = F.relu(x)
17        x = self.fc2(x)
18        x = F.relu(x)
19        x = self.fc3(x)
20        x = F.relu(x)
21        x = self.fc4(x)
22        x = F.relu(x)
23        x = self.fc5(x)
24        x = F.relu(x)
25        x = self.fc6(x)
26        x = F.relu(x)
27        x = self.fc7(x)

```

```

28         output = x
29         return output
30 fc = FC()

```

結果由下表所示，在 FC7 在 Batch Size 為 50 的情況下，效果並不是很理想，跟 FC3 的結果相近，反而是 FC4-1 跟 FC4-2 這兩者有著不錯的表現。前者 FC4-1 在 Batch Size 為 50 時有 0.954，在 Batch Size 為 30 時，FC4-1 第一次實機測試有 0.941，而第二次實機測試則有 0.930，FC4-2 則在 Batch Size 為 30 時做到 0.961。其他詳細結果放在附件 1 與附件 2。

FC 代號	Epoch	Train Loss	Test Accuracy	Batch Size	備註
FC3	1	0.2856	0.894	50	
FC3	1	0.3801	0.913	50	
FC4-1	1	0.1782	0.954	50	
FC7	1	0.2008	0.886	50	
FC4-1	1	0.2815	0.941	30	
FC4-1	1	0.0679	0.930	30	
FC4-2	1	0.0153	0.961	30	
FC7	1	0.0089	0.928	10	

在此測試只是單純的增加層，而且該測試並無做到相當全面的程度，實際上還可在投入 dropout 或者 BN 技術等，應該能夠得到更好的效果。接下來加入卷積並將 EPOCH 設為 10，其範例可以在專案目錄中的 mnist-convolutional.ipynb 找到，其改動程式碼如下：

```

1 import torch
2 import torch.nn as nn
3 import torch.utils.data as Data
4 import torchvision
5 import torch.nn.functional as F
6 import numpy as np
7
8 # torch.manual_seed(1)
9
10 EPOCH = 10
11 LR = 0.001
12 DOWNLOAD_MNIST = True
13
14 train_data = torchvision.datasets.MNIST(root='./mnist/', train=True,
15                                         transform=torchvision.transforms.ToTensor(),
16                                         download=DOWNLOAD_MNIST, )
17 test_data = torchvision.datasets.MNIST(root='./mnist/', train=False)
18
19 train_x = torch.unsqueeze(train_data.data, dim=1).type(torch.FloatTensor)
20 / 255.

```

```

20 train_y = train_data.targets
21 print(train_x.shape)
22
23 test_x = torch.unsqueeze(test_data.data, dim=1).type(torch.FloatTensor)
    [:2000] / 255. # Tensor on GPU
24 test_y = test_data.targets[:2000]
25 test_x.shape
26 import matplotlib.pyplot as plt
27 plt.imshow(test_x[1,0,:,:].numpy(), 'gray')
28 # 繪製多個 train_data
29 figure = plt.figure(figsize=(10, 8))
30 cols, rows = 5, 5
31 for i in range(1, cols * rows + 1):
32     sample_idx = torch.randint(len(train_data), size=(1,)).item()
33     img, label = train_data[sample_idx]
34     figure.add_subplot(rows, cols, i)
35     plt.title(label)
36     plt.axis("off")
37     plt.imshow(img.squeeze(), cmap="gray")
38 plt.show()
39 test_y[:10]
40 class FC(nn.Module):
41     def __init__(self):
42         super(FC, self).__init__()
43         self.conv1 = nn.Sequential(
44             nn.Conv2d(
45                 in_channels=1,
46                 out_channels=16,
47                 kernel_size=5,
48                 stride=1,
49                 padding=2,
50             ),
51             nn.ReLU(),
52             nn.MaxPool2d(kernel_size=2),
53         )
54         self.conv2 = nn.Sequential(
55             nn.Conv2d(16, 32, 5, 1, 2),
56             nn.ReLU(),
57             nn.MaxPool2d(2),
58         )
59         self.out = nn.Linear(32 * 7 * 7, 10)

```



```

60
61     def forward(self, x):
62         x = self.conv1(x)
63         x = self.conv2(x)
64         # flatten the output of conv2 to (batch_size, 32 * 7 * 7)
65         x = x.view(x.size(0), -1)
66         output = self.out(x)
67         return output
68 fc = FC()
69 optimizer = torch.optim.Adam(fc.parameters(), lr=LR)
70 #loss_func = nn.MSELoss()
71 loss_func = nn.CrossEntropyLoss()
72
73 data_size = 50000
74 batch_size = 50
75 for epoch in range(EPOCH):
76     random_indx = np.random.permutation(data_size)
77     for batch_i in range(data_size // batch_size):
78         indx = random_indx[batch_i * batch_size:(batch_i + 1) *
79                             batch_size]
79
80         b_x = train_x[indx, :]
81         b_y = train_y[indx]
82         # print(b_x.shape)
83         # print(b_y.shape)
84         output = fc(b_x)
85         loss = loss_func(output, b_y)
86         loss.backward()
87         optimizer.step()
88         optimizer.zero_grad()
89
90         if batch_i % 50 == 0:
91             test_output = fc(test_x)
92             pred_y = torch.max(test_output, 1)[1].data.squeeze()
93             # pred_y = torch.max(test_output, 1)[1].data.squeeze()
94             accuracy = torch.sum(pred_y == test_y).type(torch.FloatTensor
95                                     ) / test_y.size(0)
96             print('Epoch: ', epoch, '| train loss: %.4f' % loss.data.cpu
97                 ().numpy(), '| test accuracy: %.3f' % accuracy)
98
99 test_output = fc(test_x[:10])

```

```

98 pred_y = torch.max(test_output, 1)[1].data.squeeze() # move the
    computation in GPU
99
100 print(pred_y, 'prediction number')
101 print(test_y[:10], 'real number')
102 test_output = fc(test_x[:1])
103 pred_y = torch.max(test_output, 1)[1].data.squeeze() # move the
    computation in GPU
104
105 print(pred_y, 'prediction number')
106 print(test_y[:1], 'real number')
107
108 test_output
109 test_x[:1].shape
110 plt.imshow(test_x[:1].numpy().squeeze(), 'gray')
111 import torch
112 torch.eye(10)
113 0, 1, 2, ..., 9

```

結果由下表所示，新的版本加入卷積與改良後，在 Batch Size 為 50，EPOCH 設為 10，LR 設為 0.01，在 EPOCH 跑到 1 時，Test Accuracy 就已經到了 0.976，在 EPOCH 跑到 2 時，Test Accuracy 就已經到了 0.980，過程中 Test Accuracy 曾經到 0.989，EPOCH 跑到 9 結束時，Test Accuracy 最後到達 0.984。在此將實際結果最後一個值列在下表，同時詳細結果放在附件 3。

Epoch	Train Loss	Test Accuracy	Batch Size	備註
0	0.1125	0.973	50	
1	0.0183	0.982	50	
2	0.0788	0.982	50	
3	0.0082	0.984	50	
4	0.0563	0.984	50	
5	0.0365	0.983	50	
6	0.0039	0.986	50	
7	0.0569	0.986	50	
8	0.0004	0.989	50	
9	0.0641	0.984	50	

3 Tensorflow MNIST Training

Tensorflow MNIST Training 的 Code 為 Google 在其 Tensorflow 平臺的官方範例嘗試，而環境部屬則與 Pytorch 相似，在此則不詳述。

```

1 # Training a neural network on MNIST with Keras
2 import tensorflow as tf
3 import tensorflow_datasets as tfds
4 (ds_train, ds_test), ds_info = tfds.load(
5     'mnist',
6     split=['train', 'test'],
7     shuffle_files=True,
8     as_supervised=True,
9     with_info=True,
10 )
11
12 show = tfds.show_examples(ds_test, ds_info)
13
14 print("Number of original training examples:", len(ds_train))
15 print("Number of original test examples:", len(ds_test))
16
17 def normalize_img(image, label):
18     """Normalizes images: `uint8` -> `float32`."""
19     return tf.cast(image, tf.float32) / 255., label
20
21 ds_train = ds_train.map(
22     normalize_img, num_parallel_calls=tf.data.AUTOTUNE)
23 ds_train = ds_train.cache()
24 ds_train = ds_train.shuffle(ds_info.splits['train'].num_examples)
25 ds_train = ds_train.batch(128)
26 ds_train = ds_train.prefetch(tf.data.AUTOTUNE)
27 ds_test = ds_test.map(
28     normalize_img, num_parallel_calls=tf.data.AUTOTUNE)
29 ds_test = ds_test.batch(128)
30 ds_test = ds_test.cache()
31 ds_test = ds_test.prefetch(tf.data.AUTOTUNE)
32
33 model = tf.keras.models.Sequential([
34     tf.keras.layers.Flatten(input_shape=(28, 28)),
35     tf.keras.layers.Dense(128, activation='relu'),
36     tf.keras.layers.Dense(10)
37 ])

```

```

38 model.compile(
39     optimizer=tf.keras.optimizers.Adam(0.001),
40     loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
41     metrics=[tf.keras.metrics.SparseCategoricalAccuracy()],
42 )
43
44 model.fit(
45     ds_train,
46     epochs=6,
47     validation_data=ds_test,
48 )

```

```

In [25]: model = tf.keras.models.Sequential([
          tf.keras.layers.Flatten(input_shape=(28, 28)),
          tf.keras.layers.Dense(128, activation='relu'),
          tf.keras.layers.Dense(10)
        ])
        model.compile(
            optimizer=tf.keras.optimizers.Adam(0.001),
            loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
            metrics=[tf.keras.metrics.SparseCategoricalAccuracy()],
        )

        model.fit(
            ds_train,
            epochs=6,
            validation_data=ds_test,
        )

Epoch 1/6
469/469 [=====] - 3s 2ms/step - loss: 0.3552 - sparse_categorical_accuracy: 0.9015 - val_loss: 0.1884 - val_sparse_categorical_accuracy: 0.9479
Epoch 2/6
469/469 [=====] - 1s 2ms/step - loss: 0.1643 - sparse_categorical_accuracy: 0.9531 - val_loss: 0.1362 - val_sparse_categorical_accuracy: 0.9612
Epoch 3/6
469/469 [=====] - 1s 2ms/step - loss: 0.1181 - sparse_categorical_accuracy: 0.9663 - val_loss: 0.1064 - val_sparse_categorical_accuracy: 0.9685
Epoch 4/6
469/469 [=====] - 1s 2ms/step - loss: 0.0904 - sparse_categorical_accuracy: 0.9740 - val_loss: 0.0895 - val_sparse_categorical_accuracy: 0.9734
Epoch 5/6
469/469 [=====] - 1s 2ms/step - loss: 0.0722 - sparse_categorical_accuracy: 0.9794 - val_loss: 0.0893 - val_sparse_categorical_accuracy: 0.9720
Epoch 6/6
469/469 [=====] - 1s 2ms/step - loss: 0.0604 - sparse_categorical_accuracy: 0.9826 - val_loss: 0.0806 - val_sparse_categorical_accuracy: 0.9746
Out[25]: <keras.callbacks.History at 0x7fb69ab99c40>

```

Fig. 2. Tensorflow MNIST Training

4 Convolutional Neural Networks

在此將 CNN 歷年來重要的研究文獻整理如下所示:

方法名	年代	研究者	備註
LeNet	1998	Y. Lecun et al.	
AlexNet	2012	Alex Krizhevsky et al.	
VGG	2014.09	Karen Simonyan et al.	
Inception Net	2014.09	Christian Szegedy et al.	Google Inc.
Inception Net V2	2015.02	Sergey Ioffe et al.	Google Inc.
Inception Net V3	2015.12	Christian Szegedy et al.	Google Inc.
Inception Net V4	2016.02	Christian Szegedy et al.	Google Inc.
Xception	2016.10	François Chollet	Google Inc.
ResNet	2015.12	Kaiming He et al.	Microsoft Research
DenseNet	2016.08	Gao Huang et al.	
MobileNet V1	2017.04	Andrew G. Howard et al.	Google Inc.
MobileNet V2	2018.01	Mark Sandler et al.	Google Inc.
MobileNet V3	2019.05	Andrew Howard et al.	Google AI, Google Brain
ShuffleNet V1	2017.07	Xiangyu Zhang et al.	
ShuffleNet V2	2018.07	Xiangyu Zhang et al.	

1. Y. Lecun et al., Gradient-based learning applied to document recognition, LeNet, 1998.
2. Alex Krizhevsky et al., ImageNet Classification with Deep Convolutional Neural Networks, AlexNet, 2012.
3. Karen Simonyan et al., Very Deep Convolutional Networks for Large-Scale Image Recognition, VGG, 2014.
4. Christian Szegedy et al., Going Deeper with Convolutions, Inception Net, 2014, Google Inc.
5. Sergey Ioffe et al., Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift, Inception Net V2 , 2015, Google Inc.
6. Christian Szegedy et al., Rethinking the Inception Architecture for Computer Vision, Inception Net V3, 2015, Google Inc.
7. Christian Szegedy et al., Inception-v4, Inception-ResNet and the Impact of Residual Connections on Learning, Inception Net V4, 2016, Google Inc.
8. François Chollet, Xception: Deep Learning with Depthwise Separable Convolutions, Xception, 2016, Google Inc.
9. Kaiming He et al., Deep Residual Learning for Image Recognition, ResNet, 2015, Microsoft Research
10. Gao Huang et al., Densely Connected Convolutional Networks, DenseNet, 2016

11. Andrew G. Howard et al., MobileNets: Efficient Convolutional Neural Networks for Mobile Vision Applications, MobileNet V1, 2017, Google Inc.
12. Mark Sandler et al., MobileNetV2: Inverted Residuals and Linear Bottlenecks, MobileNet V2, 2018, Google Inc.
13. Andrew Howard et al., Searching for MobileNetV3, MobileNet V3, 2019, Google AI and Google Brain.
14. Xiangyu Zhang et al., ShuffleNet: An Extremely Efficient Convolutional Neural Network for Mobile Devices, ShuffleNet V1, 2017.
15. Xiangyu Zhang et al., ShuffleNet V2: Practical Guidelines for Efficient CNN Architecture Design, ShuffleNet V2, 2018.

5 附件 1

下為 FC4-2 則在 Batch Size 為 30 時做到 0.961 的實際測試狀況。

```
Epoch: 0 | train loss: 2.3179 | test accuracy: 0.108
Epoch: 0 | train loss: 1.7215 | test accuracy: 0.438
Epoch: 0 | train loss: 0.9065 | test accuracy: 0.595
Epoch: 0 | train loss: 0.7863 | test accuracy: 0.695
Epoch: 0 | train loss: 0.7010 | test accuracy: 0.753
Epoch: 0 | train loss: 0.2025 | test accuracy: 0.772
Epoch: 0 | train loss: 0.5316 | test accuracy: 0.784
Epoch: 0 | train loss: 0.4025 | test accuracy: 0.808
Epoch: 0 | train loss: 0.4179 | test accuracy: 0.840
Epoch: 0 | train loss: 0.2629 | test accuracy: 0.851
Epoch: 0 | train loss: 0.5766 | test accuracy: 0.885
Epoch: 0 | train loss: 0.6210 | test accuracy: 0.878
Epoch: 0 | train loss: 0.3336 | test accuracy: 0.901
Epoch: 0 | train loss: 0.1728 | test accuracy: 0.917
Epoch: 0 | train loss: 0.2423 | test accuracy: 0.919
Epoch: 0 | train loss: 0.0907 | test accuracy: 0.932
Epoch: 0 | train loss: 0.1517 | test accuracy: 0.934
Epoch: 0 | train loss: 0.5776 | test accuracy: 0.928
Epoch: 0 | train loss: 0.2446 | test accuracy: 0.924
Epoch: 0 | train loss: 0.1774 | test accuracy: 0.938
Epoch: 0 | train loss: 0.1122 | test accuracy: 0.939
Epoch: 0 | train loss: 0.5981 | test accuracy: 0.932
Epoch: 0 | train loss: 0.0891 | test accuracy: 0.939
Epoch: 0 | train loss: 0.2379 | test accuracy: 0.941
Epoch: 0 | train loss: 0.1671 | test accuracy: 0.933
Epoch: 0 | train loss: 0.1797 | test accuracy: 0.945
Epoch: 0 | train loss: 0.2858 | test accuracy: 0.941
Epoch: 0 | train loss: 0.0684 | test accuracy: 0.942
Epoch: 0 | train loss: 0.1976 | test accuracy: 0.945
Epoch: 0 | train loss: 0.2697 | test accuracy: 0.941
Epoch: 0 | train loss: 0.1319 | test accuracy: 0.938
Epoch: 0 | train loss: 0.2735 | test accuracy: 0.947
Epoch: 0 | train loss: 0.2526 | test accuracy: 0.928
Epoch: 0 | train loss: 0.0490 | test accuracy: 0.934
Epoch: 1 | train loss: 0.0266 | test accuracy: 0.945
Epoch: 1 | train loss: 0.3994 | test accuracy: 0.949
Epoch: 1 | train loss: 0.4008 | test accuracy: 0.948
Epoch: 1 | train loss: 0.1298 | test accuracy: 0.954
```

```
Epoch: 1 | train loss: 0.1138 | test accuracy: 0.946
Epoch: 1 | train loss: 0.1101 | test accuracy: 0.953
Epoch: 1 | train loss: 0.0453 | test accuracy: 0.949
Epoch: 1 | train loss: 0.0528 | test accuracy: 0.955
Epoch: 1 | train loss: 0.0294 | test accuracy: 0.950
Epoch: 1 | train loss: 0.0197 | test accuracy: 0.949
Epoch: 1 | train loss: 0.1641 | test accuracy: 0.946
Epoch: 1 | train loss: 0.1442 | test accuracy: 0.937
Epoch: 1 | train loss: 0.2131 | test accuracy: 0.951
Epoch: 1 | train loss: 0.0344 | test accuracy: 0.950
Epoch: 1 | train loss: 0.1194 | test accuracy: 0.955
Epoch: 1 | train loss: 0.0861 | test accuracy: 0.952
Epoch: 1 | train loss: 0.6072 | test accuracy: 0.957
Epoch: 1 | train loss: 0.1805 | test accuracy: 0.952
Epoch: 1 | train loss: 0.0142 | test accuracy: 0.956
Epoch: 1 | train loss: 0.0102 | test accuracy: 0.953
Epoch: 1 | train loss: 0.3421 | test accuracy: 0.955
Epoch: 1 | train loss: 0.0862 | test accuracy: 0.952
Epoch: 1 | train loss: 0.3688 | test accuracy: 0.956
Epoch: 1 | train loss: 0.0111 | test accuracy: 0.954
Epoch: 1 | train loss: 0.3614 | test accuracy: 0.955
Epoch: 1 | train loss: 0.0869 | test accuracy: 0.952
Epoch: 1 | train loss: 0.0213 | test accuracy: 0.956
Epoch: 1 | train loss: 0.2109 | test accuracy: 0.962
Epoch: 1 | train loss: 0.2889 | test accuracy: 0.951
Epoch: 1 | train loss: 0.0637 | test accuracy: 0.962
Epoch: 1 | train loss: 0.0596 | test accuracy: 0.959
Epoch: 1 | train loss: 0.0034 | test accuracy: 0.965
Epoch: 1 | train loss: 0.0190 | test accuracy: 0.963
Epoch: 1 | train loss: 0.0153 | test accuracy: 0.961
tensor([7, 2, 1, 0, 4, 1, 4, 9, 6, 9]) prediction number
tensor([7, 2, 1, 0, 4, 1, 4, 9, 5, 9]) real number
```


6 附件 2

下為 FC4-1 在 Batch Size 為 50 時有 0.954 的實際測試狀況。

```
Epoch: 0 | train loss: 2.3079 | test accuracy: 0.105
Epoch: 0 | train loss: 1.2954 | test accuracy: 0.518
Epoch: 0 | train loss: 0.8506 | test accuracy: 0.718
Epoch: 0 | train loss: 0.4927 | test accuracy: 0.789
Epoch: 0 | train loss: 0.5421 | test accuracy: 0.832
Epoch: 0 | train loss: 0.3407 | test accuracy: 0.846
Epoch: 0 | train loss: 0.5657 | test accuracy: 0.848
Epoch: 0 | train loss: 0.3150 | test accuracy: 0.881
Epoch: 0 | train loss: 0.3149 | test accuracy: 0.882
Epoch: 0 | train loss: 0.7052 | test accuracy: 0.892
Epoch: 0 | train loss: 0.2937 | test accuracy: 0.905
Epoch: 0 | train loss: 0.3780 | test accuracy: 0.909
Epoch: 0 | train loss: 0.1319 | test accuracy: 0.919
Epoch: 0 | train loss: 0.2125 | test accuracy: 0.919
Epoch: 0 | train loss: 0.4903 | test accuracy: 0.915
Epoch: 0 | train loss: 0.0368 | test accuracy: 0.925
Epoch: 0 | train loss: 0.2886 | test accuracy: 0.925
Epoch: 0 | train loss: 0.3133 | test accuracy: 0.930
Epoch: 0 | train loss: 0.1374 | test accuracy: 0.927
Epoch: 0 | train loss: 0.2062 | test accuracy: 0.939
Epoch: 1 | train loss: 0.1921 | test accuracy: 0.937
Epoch: 1 | train loss: 0.2336 | test accuracy: 0.937
Epoch: 1 | train loss: 0.2420 | test accuracy: 0.937
Epoch: 1 | train loss: 0.0554 | test accuracy: 0.942
Epoch: 1 | train loss: 0.1244 | test accuracy: 0.938
Epoch: 1 | train loss: 0.0200 | test accuracy: 0.937
Epoch: 1 | train loss: 0.1651 | test accuracy: 0.947
Epoch: 1 | train loss: 0.1534 | test accuracy: 0.942
Epoch: 1 | train loss: 0.1084 | test accuracy: 0.947
Epoch: 1 | train loss: 0.1853 | test accuracy: 0.943
Epoch: 1 | train loss: 0.0954 | test accuracy: 0.951
Epoch: 1 | train loss: 0.2550 | test accuracy: 0.951
Epoch: 1 | train loss: 0.0765 | test accuracy: 0.946
Epoch: 1 | train loss: 0.3135 | test accuracy: 0.948
Epoch: 1 | train loss: 0.0769 | test accuracy: 0.948
Epoch: 1 | train loss: 0.2409 | test accuracy: 0.948
Epoch: 1 | train loss: 0.0431 | test accuracy: 0.952
Epoch: 1 | train loss: 0.0727 | test accuracy: 0.956
```

```
Epoch: 1 | train loss: 0.0572 | test accuracy: 0.948  
Epoch: 1 | train loss: 0.1782 | test accuracy: 0.954  
tensor([7, 2, 1, 0, 4, 1, 4, 9, 5, 9]) prediction number  
tensor([7, 2, 1, 0, 4, 1, 4, 9, 5, 9]) real number
```

7 附件 3

Batch Size 為 50 ，EPOCH 為 10 ，加入卷積後改良的實際測試狀況。

```
Epoch: 0 | train loss: 2.3008 | test accuracy: 0.150
Epoch: 0 | train loss: 0.4211 | test accuracy: 0.836
Epoch: 0 | train loss: 0.1977 | test accuracy: 0.894
Epoch: 0 | train loss: 0.2739 | test accuracy: 0.895
Epoch: 0 | train loss: 0.2279 | test accuracy: 0.933
Epoch: 0 | train loss: 0.0972 | test accuracy: 0.938
Epoch: 0 | train loss: 0.2361 | test accuracy: 0.951
Epoch: 0 | train loss: 0.2102 | test accuracy: 0.957
Epoch: 0 | train loss: 0.0642 | test accuracy: 0.956
Epoch: 0 | train loss: 0.0603 | test accuracy: 0.951
Epoch: 0 | train loss: 0.1301 | test accuracy: 0.965
Epoch: 0 | train loss: 0.1806 | test accuracy: 0.959
Epoch: 0 | train loss: 0.1146 | test accuracy: 0.961
Epoch: 0 | train loss: 0.0452 | test accuracy: 0.965
Epoch: 0 | train loss: 0.0467 | test accuracy: 0.967
Epoch: 0 | train loss: 0.0450 | test accuracy: 0.971
Epoch: 0 | train loss: 0.1723 | test accuracy: 0.969
Epoch: 0 | train loss: 0.2110 | test accuracy: 0.969
Epoch: 0 | train loss: 0.1457 | test accuracy: 0.970
Epoch: 0 | train loss: 0.1125 | test accuracy: 0.973
Epoch: 1 | train loss: 0.1761 | test accuracy: 0.976
Epoch: 1 | train loss: 0.1053 | test accuracy: 0.971
Epoch: 1 | train loss: 0.0772 | test accuracy: 0.976
Epoch: 1 | train loss: 0.0253 | test accuracy: 0.975
Epoch: 1 | train loss: 0.0090 | test accuracy: 0.975
Epoch: 1 | train loss: 0.0401 | test accuracy: 0.976
Epoch: 1 | train loss: 0.0589 | test accuracy: 0.978
Epoch: 1 | train loss: 0.0169 | test accuracy: 0.976
Epoch: 1 | train loss: 0.0126 | test accuracy: 0.975
Epoch: 1 | train loss: 0.0848 | test accuracy: 0.977
Epoch: 1 | train loss: 0.0188 | test accuracy: 0.979
Epoch: 1 | train loss: 0.0560 | test accuracy: 0.980
Epoch: 1 | train loss: 0.0363 | test accuracy: 0.975
Epoch: 1 | train loss: 0.0170 | test accuracy: 0.982
Epoch: 1 | train loss: 0.0317 | test accuracy: 0.975
Epoch: 1 | train loss: 0.0755 | test accuracy: 0.983
Epoch: 1 | train loss: 0.0850 | test accuracy: 0.980
```

```
Epoch: 1 | train loss: 0.0484 | test accuracy: 0.982
Epoch: 1 | train loss: 0.0240 | test accuracy: 0.982
Epoch: 1 | train loss: 0.0183 | test accuracy: 0.982
Epoch: 2 | train loss: 0.0272 | test accuracy: 0.980
Epoch: 2 | train loss: 0.0021 | test accuracy: 0.983
Epoch: 2 | train loss: 0.0064 | test accuracy: 0.981
Epoch: 2 | train loss: 0.0182 | test accuracy: 0.984
Epoch: 2 | train loss: 0.0115 | test accuracy: 0.984
Epoch: 2 | train loss: 0.0075 | test accuracy: 0.979
Epoch: 2 | train loss: 0.0054 | test accuracy: 0.985
Epoch: 2 | train loss: 0.0064 | test accuracy: 0.982
Epoch: 2 | train loss: 0.0045 | test accuracy: 0.983
Epoch: 2 | train loss: 0.0987 | test accuracy: 0.981
Epoch: 2 | train loss: 0.0129 | test accuracy: 0.984
Epoch: 2 | train loss: 0.0289 | test accuracy: 0.983
Epoch: 2 | train loss: 0.0282 | test accuracy: 0.980
Epoch: 2 | train loss: 0.0262 | test accuracy: 0.984
Epoch: 2 | train loss: 0.0643 | test accuracy: 0.976
Epoch: 2 | train loss: 0.0618 | test accuracy: 0.984
Epoch: 2 | train loss: 0.0079 | test accuracy: 0.987
Epoch: 2 | train loss: 0.0912 | test accuracy: 0.984
Epoch: 2 | train loss: 0.0409 | test accuracy: 0.984
Epoch: 2 | train loss: 0.0788 | test accuracy: 0.982
Epoch: 3 | train loss: 0.0182 | test accuracy: 0.980
Epoch: 3 | train loss: 0.0124 | test accuracy: 0.984
Epoch: 3 | train loss: 0.0248 | test accuracy: 0.984
Epoch: 3 | train loss: 0.0077 | test accuracy: 0.984
Epoch: 3 | train loss: 0.0038 | test accuracy: 0.982
Epoch: 3 | train loss: 0.0444 | test accuracy: 0.988
Epoch: 3 | train loss: 0.1470 | test accuracy: 0.983
Epoch: 3 | train loss: 0.1792 | test accuracy: 0.985
Epoch: 3 | train loss: 0.0022 | test accuracy: 0.983
Epoch: 3 | train loss: 0.0345 | test accuracy: 0.986
Epoch: 3 | train loss: 0.0163 | test accuracy: 0.981
Epoch: 3 | train loss: 0.0084 | test accuracy: 0.984
Epoch: 3 | train loss: 0.0118 | test accuracy: 0.987
Epoch: 3 | train loss: 0.0162 | test accuracy: 0.985
Epoch: 3 | train loss: 0.0096 | test accuracy: 0.984
Epoch: 3 | train loss: 0.0072 | test accuracy: 0.980
Epoch: 3 | train loss: 0.0026 | test accuracy: 0.987
Epoch: 3 | train loss: 0.0025 | test accuracy: 0.987
```

```
Epoch: 3 | train loss: 0.1035 | test accuracy: 0.984
Epoch: 3 | train loss: 0.0082 | test accuracy: 0.984
Epoch: 4 | train loss: 0.0066 | test accuracy: 0.984
Epoch: 4 | train loss: 0.0099 | test accuracy: 0.987
Epoch: 4 | train loss: 0.0360 | test accuracy: 0.987
Epoch: 4 | train loss: 0.1021 | test accuracy: 0.981
Epoch: 4 | train loss: 0.0202 | test accuracy: 0.987
Epoch: 4 | train loss: 0.0013 | test accuracy: 0.984
Epoch: 4 | train loss: 0.0011 | test accuracy: 0.983
Epoch: 4 | train loss: 0.0125 | test accuracy: 0.986
Epoch: 4 | train loss: 0.0004 | test accuracy: 0.987
Epoch: 4 | train loss: 0.0045 | test accuracy: 0.983
Epoch: 4 | train loss: 0.0026 | test accuracy: 0.987
Epoch: 4 | train loss: 0.0232 | test accuracy: 0.981
Epoch: 4 | train loss: 0.0032 | test accuracy: 0.987
Epoch: 4 | train loss: 0.0013 | test accuracy: 0.988
Epoch: 4 | train loss: 0.0003 | test accuracy: 0.985
Epoch: 4 | train loss: 0.0051 | test accuracy: 0.984
Epoch: 4 | train loss: 0.0044 | test accuracy: 0.983
Epoch: 4 | train loss: 0.0118 | test accuracy: 0.985
Epoch: 4 | train loss: 0.0135 | test accuracy: 0.984
Epoch: 4 | train loss: 0.0563 | test accuracy: 0.984
Epoch: 5 | train loss: 0.0236 | test accuracy: 0.988
Epoch: 5 | train loss: 0.0005 | test accuracy: 0.987
Epoch: 5 | train loss: 0.0054 | test accuracy: 0.983
Epoch: 5 | train loss: 0.0042 | test accuracy: 0.989
Epoch: 5 | train loss: 0.0013 | test accuracy: 0.985
Epoch: 5 | train loss: 0.0023 | test accuracy: 0.985
Epoch: 5 | train loss: 0.0003 | test accuracy: 0.985
Epoch: 5 | train loss: 0.0038 | test accuracy: 0.982
Epoch: 5 | train loss: 0.0113 | test accuracy: 0.987
Epoch: 5 | train loss: 0.0019 | test accuracy: 0.985
Epoch: 5 | train loss: 0.0019 | test accuracy: 0.987
Epoch: 5 | train loss: 0.0026 | test accuracy: 0.985
Epoch: 5 | train loss: 0.0415 | test accuracy: 0.987
Epoch: 5 | train loss: 0.0559 | test accuracy: 0.986
Epoch: 5 | train loss: 0.0125 | test accuracy: 0.987
Epoch: 5 | train loss: 0.0121 | test accuracy: 0.984
Epoch: 5 | train loss: 0.0029 | test accuracy: 0.984
Epoch: 5 | train loss: 0.0102 | test accuracy: 0.984
Epoch: 5 | train loss: 0.0012 | test accuracy: 0.987
```

Epoch: 5 | train loss: 0.0365 | test accuracy: 0.983
Epoch: 6 | train loss: 0.0462 | test accuracy: 0.987
Epoch: 6 | train loss: 0.0074 | test accuracy: 0.984
Epoch: 6 | train loss: 0.0009 | test accuracy: 0.984
Epoch: 6 | train loss: 0.0004 | test accuracy: 0.984
Epoch: 6 | train loss: 0.0428 | test accuracy: 0.987
Epoch: 6 | train loss: 0.0032 | test accuracy: 0.989
Epoch: 6 | train loss: 0.0090 | test accuracy: 0.987
Epoch: 6 | train loss: 0.0009 | test accuracy: 0.985
Epoch: 6 | train loss: 0.0626 | test accuracy: 0.985
Epoch: 6 | train loss: 0.0674 | test accuracy: 0.987
Epoch: 6 | train loss: 0.0045 | test accuracy: 0.985
Epoch: 6 | train loss: 0.0087 | test accuracy: 0.988
Epoch: 6 | train loss: 0.0381 | test accuracy: 0.987
Epoch: 6 | train loss: 0.0868 | test accuracy: 0.985
Epoch: 6 | train loss: 0.0900 | test accuracy: 0.986
Epoch: 6 | train loss: 0.0019 | test accuracy: 0.984
Epoch: 6 | train loss: 0.0713 | test accuracy: 0.983
Epoch: 6 | train loss: 0.0020 | test accuracy: 0.984
Epoch: 6 | train loss: 0.0033 | test accuracy: 0.984
Epoch: 6 | train loss: 0.0039 | test accuracy: 0.986
Epoch: 7 | train loss: 0.0912 | test accuracy: 0.988
Epoch: 7 | train loss: 0.0018 | test accuracy: 0.985
Epoch: 7 | train loss: 0.0101 | test accuracy: 0.987
Epoch: 7 | train loss: 0.0012 | test accuracy: 0.985
Epoch: 7 | train loss: 0.0204 | test accuracy: 0.985
Epoch: 7 | train loss: 0.0012 | test accuracy: 0.987
Epoch: 7 | train loss: 0.0005 | test accuracy: 0.985
Epoch: 7 | train loss: 0.0208 | test accuracy: 0.985
Epoch: 7 | train loss: 0.0024 | test accuracy: 0.987
Epoch: 7 | train loss: 0.0509 | test accuracy: 0.987
Epoch: 7 | train loss: 0.0063 | test accuracy: 0.978
Epoch: 7 | train loss: 0.0007 | test accuracy: 0.988
Epoch: 7 | train loss: 0.0007 | test accuracy: 0.985
Epoch: 7 | train loss: 0.0056 | test accuracy: 0.986
Epoch: 7 | train loss: 0.0029 | test accuracy: 0.986
Epoch: 7 | train loss: 0.0469 | test accuracy: 0.988
Epoch: 7 | train loss: 0.0042 | test accuracy: 0.985
Epoch: 7 | train loss: 0.0139 | test accuracy: 0.986
Epoch: 7 | train loss: 0.0209 | test accuracy: 0.987
Epoch: 7 | train loss: 0.0569 | test accuracy: 0.986

```
Epoch: 8 | train loss: 0.0010 | test accuracy: 0.984
Epoch: 8 | train loss: 0.0007 | test accuracy: 0.988
Epoch: 8 | train loss: 0.0288 | test accuracy: 0.988
Epoch: 8 | train loss: 0.0009 | test accuracy: 0.985
Epoch: 8 | train loss: 0.0052 | test accuracy: 0.987
Epoch: 8 | train loss: 0.0008 | test accuracy: 0.988
Epoch: 8 | train loss: 0.0005 | test accuracy: 0.985
Epoch: 8 | train loss: 0.0049 | test accuracy: 0.987
Epoch: 8 | train loss: 0.0037 | test accuracy: 0.986
Epoch: 8 | train loss: 0.0060 | test accuracy: 0.984
Epoch: 8 | train loss: 0.0019 | test accuracy: 0.984
Epoch: 8 | train loss: 0.0017 | test accuracy: 0.986
Epoch: 8 | train loss: 0.0137 | test accuracy: 0.987
Epoch: 8 | train loss: 0.0018 | test accuracy: 0.988
Epoch: 8 | train loss: 0.0003 | test accuracy: 0.989
Epoch: 8 | train loss: 0.0186 | test accuracy: 0.985
Epoch: 8 | train loss: 0.0170 | test accuracy: 0.985
Epoch: 8 | train loss: 0.0005 | test accuracy: 0.980
Epoch: 8 | train loss: 0.0109 | test accuracy: 0.988
Epoch: 8 | train loss: 0.0004 | test accuracy: 0.989
Epoch: 9 | train loss: 0.0152 | test accuracy: 0.987
Epoch: 9 | train loss: 0.0130 | test accuracy: 0.987
Epoch: 9 | train loss: 0.0066 | test accuracy: 0.987
Epoch: 9 | train loss: 0.0027 | test accuracy: 0.988
Epoch: 9 | train loss: 0.0043 | test accuracy: 0.986
Epoch: 9 | train loss: 0.0088 | test accuracy: 0.987
Epoch: 9 | train loss: 0.0025 | test accuracy: 0.988
Epoch: 9 | train loss: 0.0029 | test accuracy: 0.988
Epoch: 9 | train loss: 0.0042 | test accuracy: 0.987
Epoch: 9 | train loss: 0.0001 | test accuracy: 0.987
Epoch: 9 | train loss: 0.0005 | test accuracy: 0.988
Epoch: 9 | train loss: 0.0038 | test accuracy: 0.985
Epoch: 9 | train loss: 0.0004 | test accuracy: 0.984
Epoch: 9 | train loss: 0.0063 | test accuracy: 0.985
Epoch: 9 | train loss: 0.0067 | test accuracy: 0.987
Epoch: 9 | train loss: 0.0169 | test accuracy: 0.984
Epoch: 9 | train loss: 0.0034 | test accuracy: 0.985
Epoch: 9 | train loss: 0.0027 | test accuracy: 0.985
Epoch: 9 | train loss: 0.0007 | test accuracy: 0.983
Epoch: 9 | train loss: 0.0641 | test accuracy: 0.984
tensor([7, 2, 1, 0, 4, 1, 4, 9, 5, 9]) prediction number
```

```
tensor([7, 2, 1, 0, 4, 1, 4, 9, 5, 9]) real number
```